


File Systems



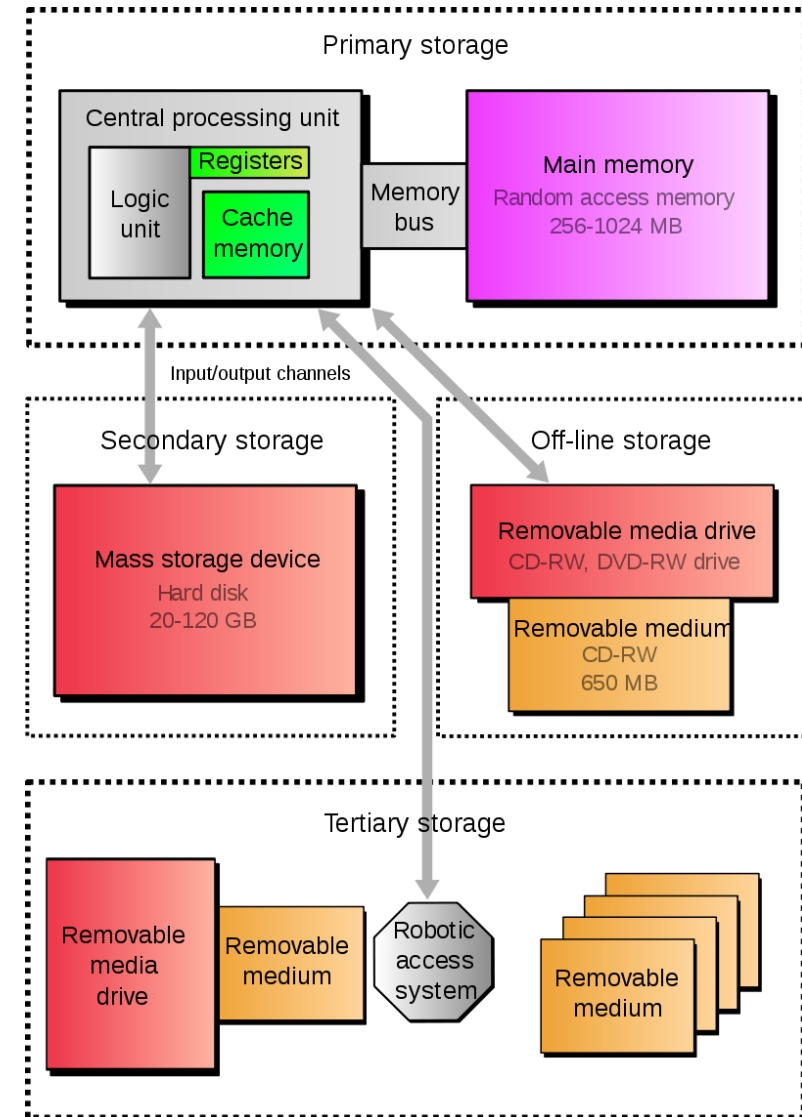
File System

- A file system or filesystem (often abbreviated to fs) is a method and data structure that the operating system uses to control how data is stored and retrieved.



Data Storage forms

- **Primary storage** (also known as main memory, internal memory, or prime memory), often referred to simply as memory, is the only one directly accessible to the CPU.
- **Secondary storage** (also known as external memory or auxiliary storage) differs from primary storage in that it is not directly accessible by the CPU.
- **Tertiary storage** or tertiary memory is a level below secondary storage. Typically, it involves a robotic mechanism which will mount (insert) and dismount removable mass storage media into a storage device according to the system's demands; such data are often copied to secondary storage before use.





Hard disk drive

- A hard disk drive (HDD) is an electro-mechanical data storage device that stores and retrieves digital data using magnetic storage with one or more rigid rapidly rotating platters coated with magnetic material.
- Introduced by IBM in 1956.



HDD History Examples

- IBM 350 (First Hard disk)
 - Introduced in 1956.
 - Size: two large refrigerators.
 - Capacity: 3.75 Mbyte.



HDD History Examples

- IBM 1311
 - Introduced in 1962.
 - Size: washing machine.
 - Capacity: 1.43 Mbyte.




HDD Improvements over years


Improvement of HDD characteristics over time

Parameter	Started with (1957)	Improved to	Improvement
Capacity (formatted)	3.75 megabytes ^[18]	22 terabytes (as of 2023) ^[19]	5.86-million-to-one ^[c]
Physical volume	68 cubic feet (1.9 m ³) ^{[d][6]}	2.1 cubic inches (34 cm ³) ^{[20][e]}	56,000-to-one ^[f]
Weight	2,000 pounds (910 kg) ^[6]	2.2 ounces (62 g) ^[20]	15,000-to-one ^[g]
Average access time	approx. 600 milliseconds ^[6]	2.5 ms to 10 ms; RW RAM dependent	about 200-to-one ^[h]
Price	US\$9,200 per megabyte (1961; US\$83,107 in 2021) ^[21]	US\$0.024 per gigabyte by 2020 ^{[22][i][23]}	3.46-billion-to-one ^[i]
Data density	2,000 bits per square inch ^[24]	1.3 terabits per square inch in 2015 ^[25]	650-million-to-one ^[k]
Average lifespan	c. 2000 hrs MTBF ^[citation needed]	c. 2,500,000 hrs (~285 years) MTBF ^[26]	1250-to-one ^[j]

Components of HDD




Components of HDD




Components of HDD

- Zone bit recording (ZBR) is a method used by disk drives to optimize the tracks for increased data capacity.
- Zone recording was pioneered and patented by *Chuck Peddle* in 1961 while working for General Electric.



HDD Addressing Schemes

- CHS (Cylinder – Head – Sector)
 - Limited by:
 - 512 bytes / sector.
 - 63 sectors / track.
 - 255 heads (tracks/Cylinder).
 - 1024 Cylinders.
 - Maximum disk size: 8064 MB.
- LBA (Logical Block Addressing)



SSD

- A solid-state drive (SSD)
 - is a solid-state storage device that uses integrated circuit assemblies to store data persistently, typically using flash memory, and functioning as secondary storage in the hierarchy of computer storage.
- Outperforms the HDD in:
 - Speed.
 - Noise.
 - Capacity.
 - Durability.
 - Size.
 - Fragmentation.
 - **Price.**



Disk Partitioning

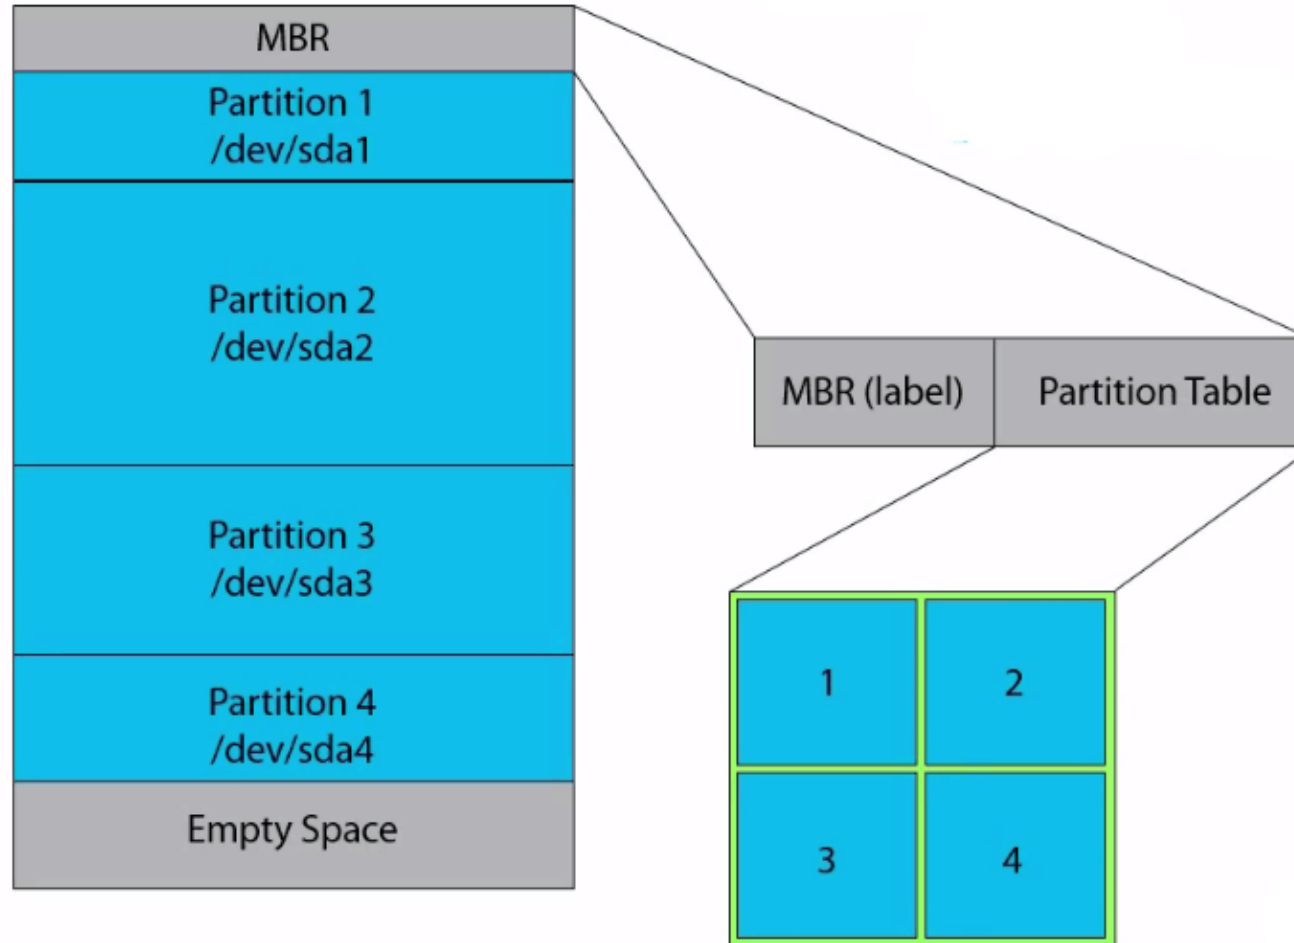
- **Disk partitioning** is the creation of one or more regions on secondary storage, so that each region can be managed separately. These regions are called **partitions**.
- Disk Partitioning reasons:
 - Organize data.
 - Security.
 - Avoid crossing size limits.
 - Implement quotas.
 - Supporting multiple filesystems.
- Partition Types:
 - Data Partitions.
 - Swap Partitions.



Disk Partitioning

- **Partition Table:** is a table maintained on a disk by the operating system that outlines and describes the partitions on that disk.
- Partition Table Types:
 - MBR (Master Boot Record).
 - GPT (GUID (Globally Unique Identifier) Partition Table).


Master Boot Record



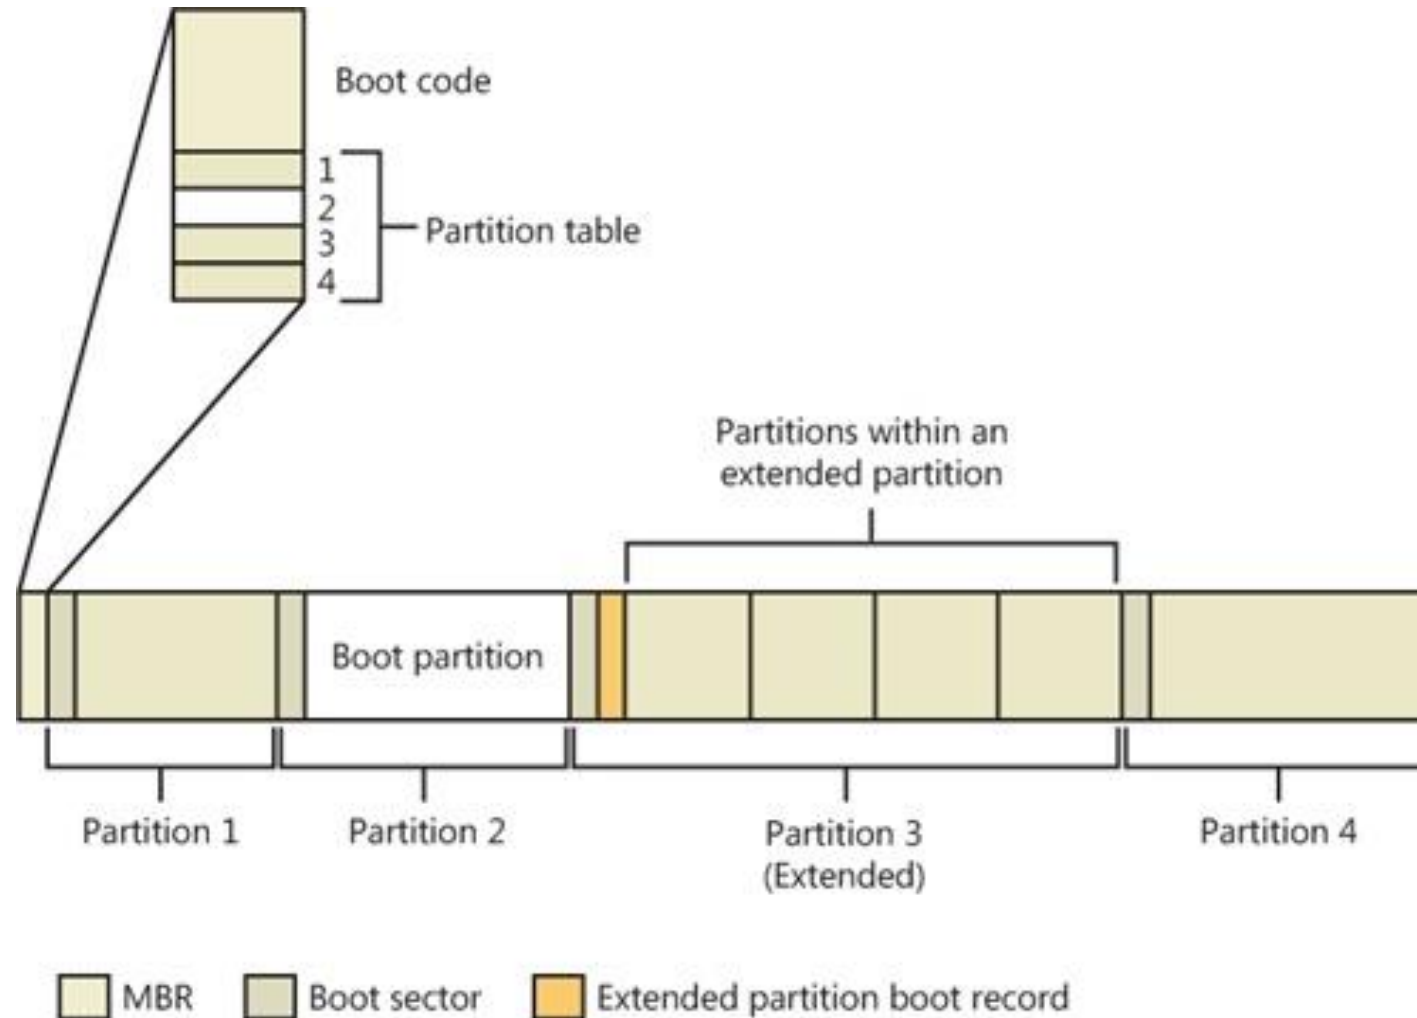
Master Boot Record



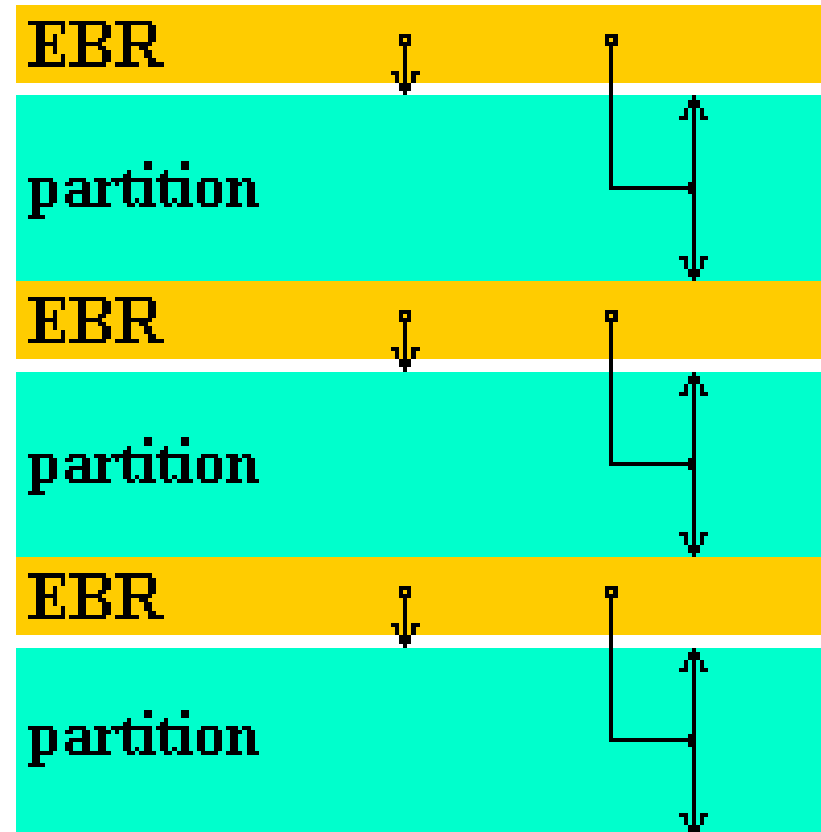
Master Boot Record




MBR with an Extended Partition



Extended Boot Record (First entry)



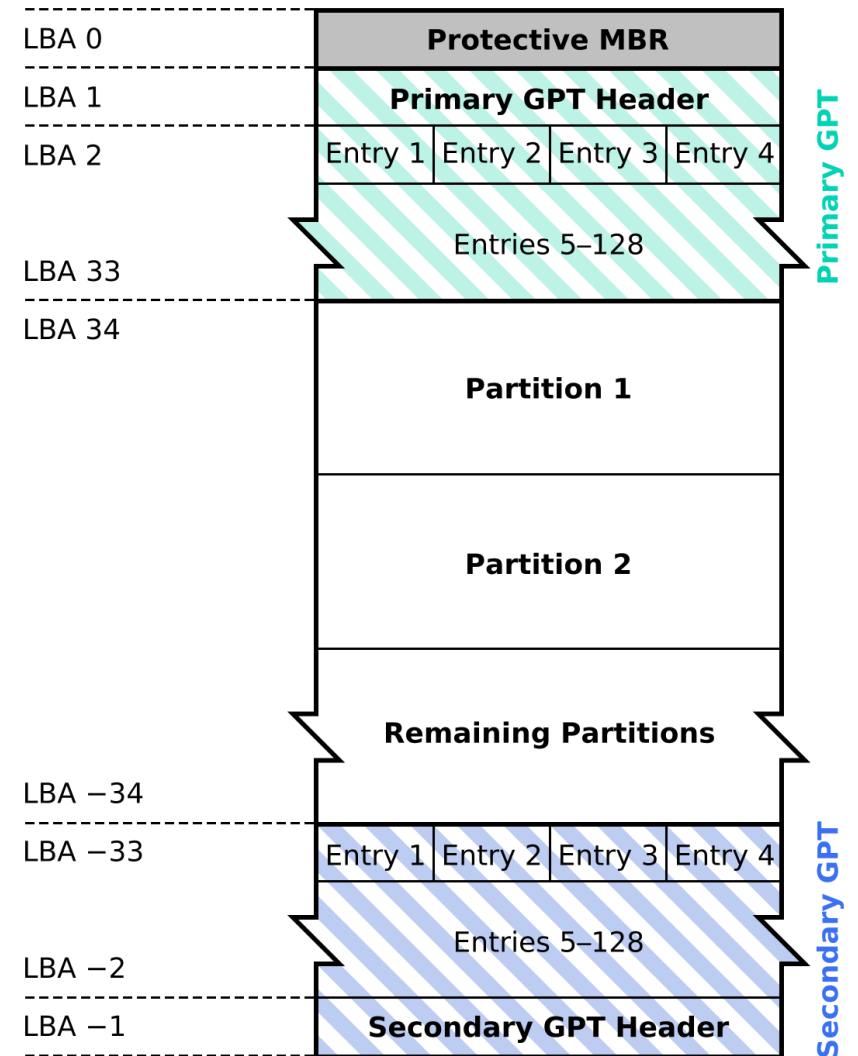
Extended Boot Record (Second entry)



GUID Partition Table (GPT)

- Up to 128 partition.
- Partition size: up to 9.44 ZiB (10^9 TiB).
- GPT table is repeated at the end of the hard disk.

GUID Partition Table Scheme



GPT Header

GPT header format

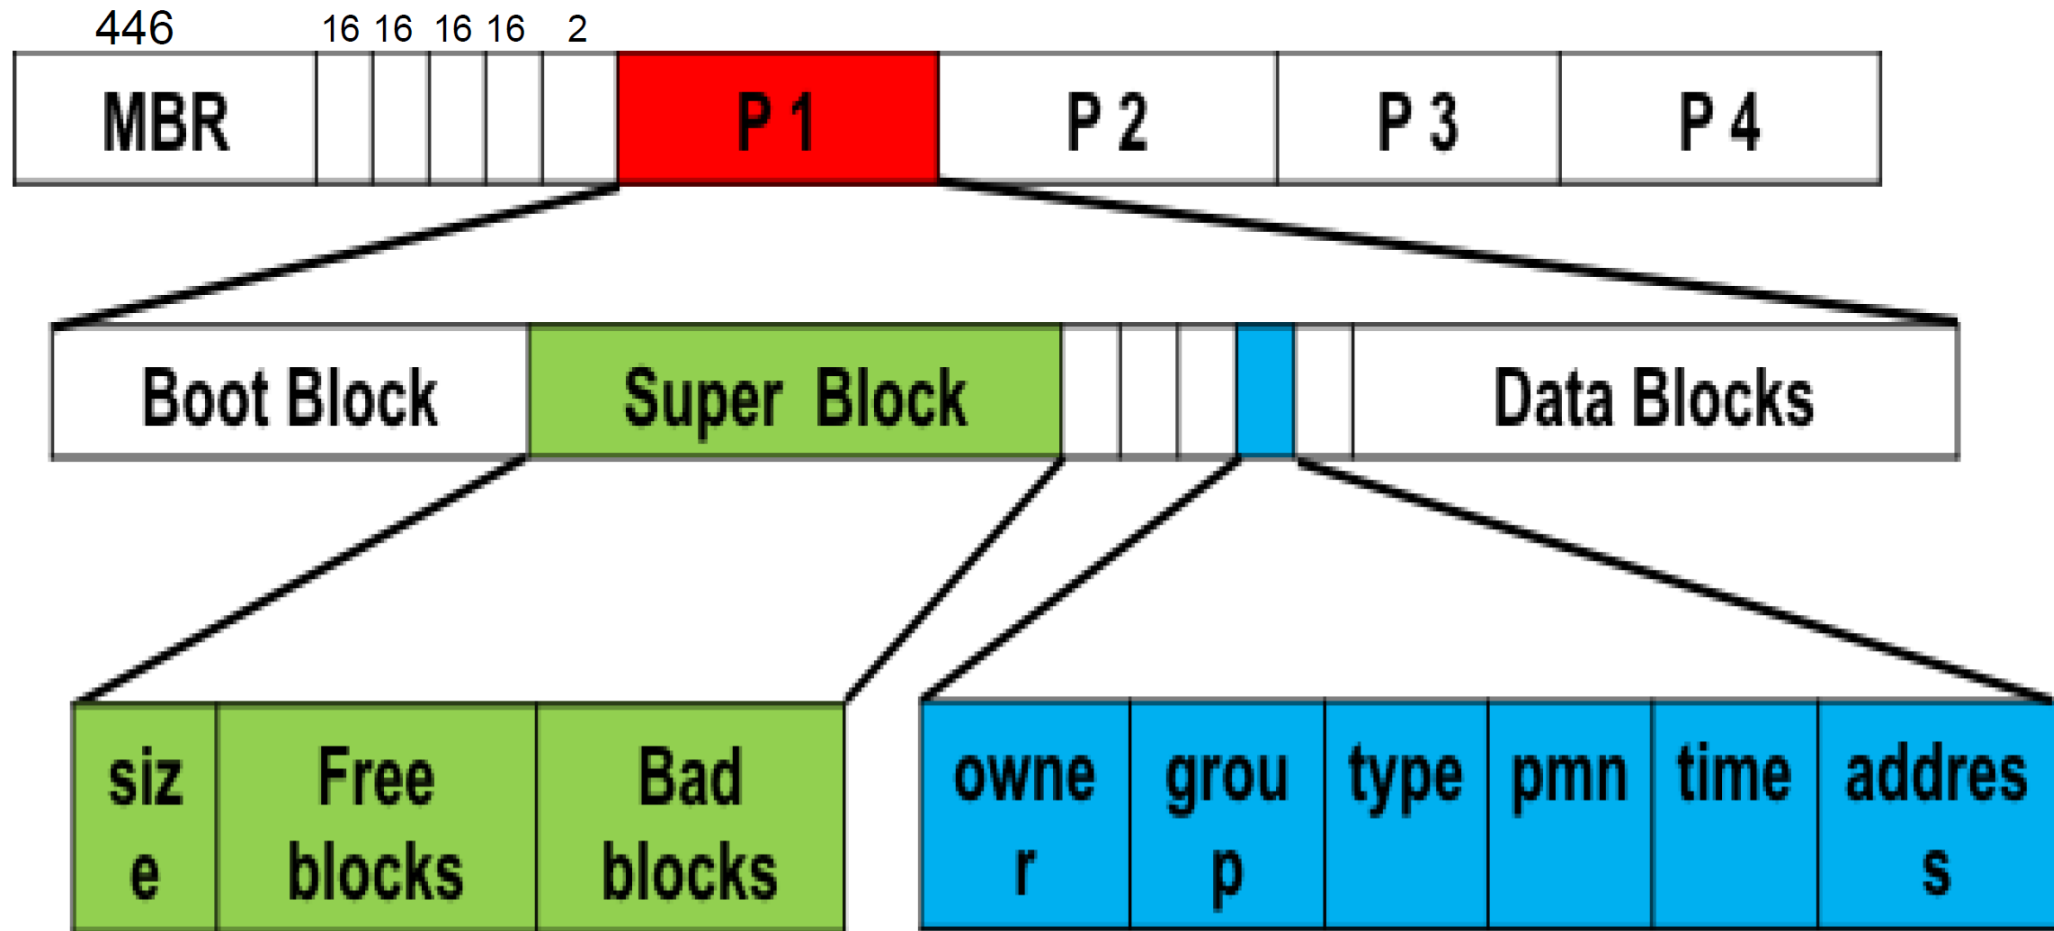
Offset	Length	Contents
0 (0x00)	8 bytes	Signature ("EFI PART", 45h 46h 49h 20h 50h 41h 52h 54h or 0x5452415020494645ULL ^[a] on little-endian machines)
8 (0x08)	4 bytes	Revision number of header - 1.0 (00h 00h 01h 00h) for UEFI 2.10
12 (0x0C)	4 bytes	Header size in little endian (in bytes, usually 5Ch 00h 00h 00h or 92 bytes)
16 (0x10)	4 bytes	CRC32 of header (offset +0 to +0x5b) in little endian, with this field zeroed during calculation
20 (0x14)	4 bytes	Reserved; must be zero
24 (0x18)	8 bytes	Current LBA (location of this header copy)
32 (0x20)	8 bytes	Backup LBA (location of the other header copy)
40 (0x28)	8 bytes	First usable LBA for partitions (primary partition table last LBA + 1)
48 (0x30)	8 bytes	Last usable LBA (secondary partition table first LBA - 1)
56 (0x38)	16 bytes	Disk GUID in mixed endian ^[12]
72 (0x48)	8 bytes	Starting LBA of array of partition entries (usually 2 for compatibility)
80 (0x50)	4 bytes	Number of partition entries in array
84 (0x54)	4 bytes	Size of a single partition entry (usually 80h or 128)
88 (0x58)	4 bytes	CRC32 of partition entries array in little endian
92 (0x5C)	*	Reserved; must be zeroes for the rest of the block (420 bytes for a sector size of 512 bytes; but can be more with larger sector sizes)

GPT Partition Entry

GUID partition entry format

Offset	Length	Contents
0 (0x00)	16 bytes	Partition type GUID (mixed endian ^[12])
16 (0x10)	16 bytes	Unique partition GUID (mixed endian)
32 (0x20)	8 bytes	First LBA (little endian)
40 (0x28)	8 bytes	Last LBA (inclusive, usually odd)
48 (0x30)	8 bytes	Attribute flags (e.g. bit 60 denotes read-only)
56 (0x38)	72 bytes	Partition name (36 UTF-16LE code units)

Unix Filesystem



Inode structure

